

Enabling naturalistic, long-duration and continual animal experimentation

March 16, 2025

1 Vision

For over four years, at the Sainsbury Wellcome Centre and Gatsby Computational Neuroscience Unit, we have been developing the AEON platform, a set of hardware and software tools that support a new type of experimentation, where animals are allowed to express ethologically-relevant behaviors, in naturalistic environments, and in long-duration experiments, while their behavior and neural activity is monitored continuously for weeks to months. We have used this platform to characterize foraging behavior in both solitary and groups of mice (?) (Figure ??).

Our US partner, the Allen Institute for Neural Dynamics, is using the AEON platform in continuous learning experiments, where mice freely explore odors continuously for days to weeks (?).

This is an unprecedented type of experimentation that ...

Several groups around the world are performing this new type of experimentation .

Aim 1: create software for visual exploration and statistical data analysis of NaLoDuCo experimental data on the cloud

Using the AEON platform we and others have recorded unprecedented data. We next want to openly disseminate these data. However, this dissemination is not trivial, as datasets generated by this new type of experimentation are

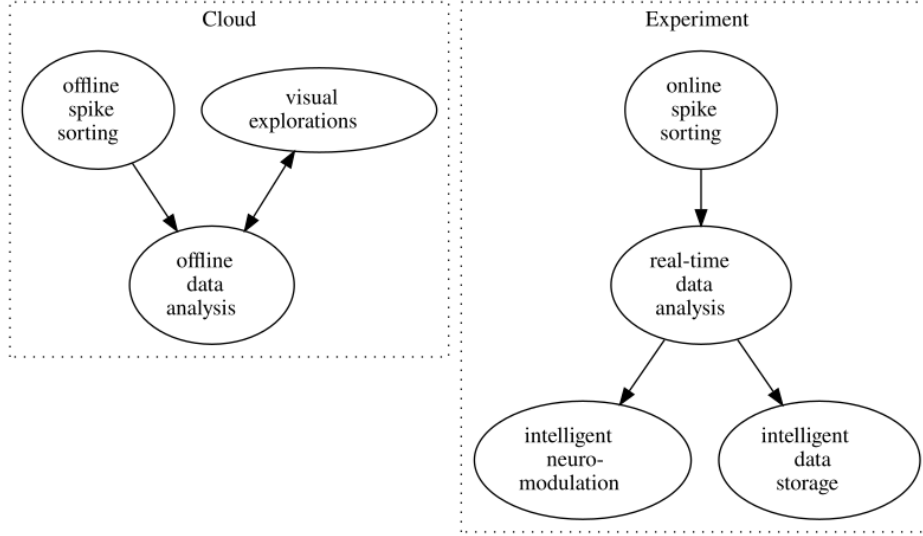


Figure 1: Proposal aims

enormous. For instance, the size of a dataset generated from a one week recording of behavioral and neural activity from a foraging mouse in SWC experiments exceeds 200 terabytes. It will take users several days to download these datasets over standard Internet connections. So, instead of bringing data to users, we will bring users to data, by storing datasets in the cloud, and providing **cloud software to allow users to visually explore and statistically analyse behavioural and neural NaLoDuCo datasets where they live** (Figure 1, left box).

Our statistical analysis of neural time series will require knowledge of the spiking activity of single units; i.e., spike sorting. In long-duration experiments with freely moving animals spike sorting is a challenging problem, because movements of recording probes change the shape of spike waveforms over time and complicate the assignment of spikes to units based on spike waveforms. We will address this problem by developing **spike sorting methods for long-duration and continual, long-duration and high-channel-count recordings** (Figure 1, left box).

Aim 2: create real-time machine learning methods for intelligent experimentation

In small-animal Neuroscience, most often statistical processing of neural time series is performed offline; i.e., experimental data is collected, saved to files, which are later statistically processed, with no runtime constraints. Most often all experimental data is processed at the same time; i.e., batch processing.

A new online statistical processing approach is now emerging in small-animal Neuroscience, where data is processed while it is being collected, and at the speed of data generation (?).

Online methods are well suited for NaLoDuCo experimentation. In experiments extending for weeks to months animals learn and adapt, their motivation and fatigue may fluctuate, and experimental conditions (e.g., lighting) may change. Offline batch processing algorithms cannot model this type of changing data. They assume stationary data whose statistical properties do not change across time. Differently, most online processing algorithms are robust to these changes. Also, NaLoDuCo experimentation is well suited for online methods, as the long-duration of these experiments provide a large amount of data to accurately fit expressive online methods.

We will **optimize methods developed for Aim 1 so that they can operate in real time**, and focus on the following two applications of these online methods (Figure 1, right box).

Intelligent neuromodulation

Brain activity can be modulated optically, chemically and electrically (). Most commonly this modulation is done at fixed experimental times, or based on simple behavioral or neural observations.

We will guide optogenetic manipulations based on inferences from advanced machine learning methods. For example, a scientist may hypothesize that a peak in a neural latent variable, inferred from a prefrontal cortex population, signals the moment when mice decide to begin a foraging bout. To test this, she runs an online machine learning model to estimate latent variables from prefrontal cortex activity, predicting when this peak will occur. She then optogenetically inactivates the neural population at the forecasted time. Because inactivation prevented the mouse from initiating a foraging bout, her hypothesis is supported.

Intelligent experimental data storage

As the duration of NaLoDuCo experiments become longer, and the richness of the behavioral and neural recordings become larger, it will be unfeasible to store all raw data. We will be forced to intelligently decide, in real time, subsets of data to discard.

For instance, if we are recording videos from a mouse foraging in a large arena with ten high-resolution cameras, it would save considerable storage if at any time we only save videos from cameras capturing the mouse at that time. This could be done by tracking the position of the mouse in real time with probabilistic machine learning methods. Then, when the confidence of the tracking is high, we would only save videos of cameras capturing the mouse at the tracked position, but when the confidence is low, we would save all videos.