

# Enabling naturalistic, long-duration and continual animal experimentation

March 24, 2025

## Contents

<b>1</b>	<b>Vision</b>	<b>2</b>
<b>2</b>	<b>Approach</b>	<b>5</b>
2.1	Offline analysis methods . . . . .	5
2.2	Initial list of methods . . . . .	6
2.2.1	Challenges . . . . .	7
2.2.2	Opportunities . . . . .	9
2.2.3	Related research . . . . .	11
2.2.4	Outputs . . . . .	12
2.3	Visual Exploration . . . . .	12
2.3.1	Outputs . . . . .	12
2.4	Spike Sorting . . . . .	12
2.4.1	Outputs . . . . .	12
2.5	Dissemination . . . . .	13
2.5.1	Outputs . . . . .	13
2.6	Online Machine Learning . . . . .	13
2.6.1	Outputs . . . . .	13
2.7	Software and Infrastructure . . . . .	13

# 1 Vision

For over four years, at the Sainsbury Wellcome Centre and Gatsby Computational Neuroscience Unit, we have been developing the AEON platform, a set of hardware and software tools that support a new type of experimentation, where animals are allowed to express ethologically-relevant behaviors, in naturalistic environments, and in long-duration experiments, while their behavior and neural activity is monitored continuously for weeks to months. We have used this platform to characterize foraging behavior in both solitary and groups of mice (?) (Figure ??).

Our US partner, the Allen Institute for Neural Dynamics, is using the AEON platform in continuous learning experiments, where mice freely explore odors continuously for days to weeks (?).

This is an unprecedented type of experimentation that ...

Several groups around the world are performing this new type of experimentation .

We have built the AEON platform, and have used it to collect weeks-to months-long NaLoDuCo experimental data. We next propose to develop advanced machine learning methods and intelligent visualisations to extract meaning from this data (Aim 1).

A central aim of both the SWC/GCNU and AIND is to contribute to open science. We thus propose to create software infrastructure to openly disseminate NaLoDuCo recordings, visualisation and data analysis methods (Aim 2).

Over more than four years we have developed the AEON platform following high software engineering practices. It is an open source platform that anybody can use and modify (?). We want it to become the standard platform for the collection of NaLoDuCo experimental data. We are currently using AEON on two new NaLoDuCo experiments: (1) odor learning experiments, lasting for days to weeks, lead by Dr. Carl Schoonover at the AIND, and (2) foraging experiments in very large arenas (eight meters in diameter), lead by Prof. Tiago Branco at the SWC. We will extend and validate the functionality of the AEON platform by applying it to these and new NaLoDuCo experiments. A key functionality that we propose to add as part of this project is real-time machine learning, to allow to control AEON experiments with live inferences (Aim 3).

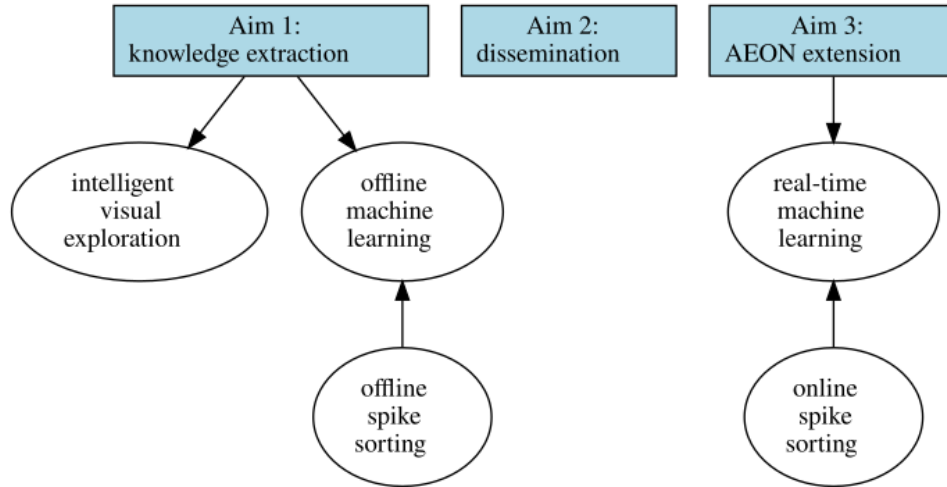


Figure 1: Proposal aims

## Aim 1: create infrastructure for open dissemination of NaLoDuCo experimental recordings

The dissemination of NaLoDuCo recordings is not trivial, as datasets generated by this new type of experimentation are enormous. For instance, the size of a dataset generated from a one week recording of behavioral and neural activity from a foraging mouse in SWC experiments exceeds 200 terabytes. It will take users several days to download these datasets over standard Internet connections.

Instead of bringing data to users, we will bring users to data, by storing datasets in the cloud (or in institutional clusters), and providing **cloud software to allow users to visually explore and statistically analyse behavioural and neural NaLoDuCo datasets where they live** (Figure 1, left box).

Our statistical analysis of neural time series will require knowledge of the spiking activity of single units; i.e., spike sorting. In long-duration experiments with freely moving animals spike sorting is a challenging problem, because movements of recording probes change the shape of spike waveforms over time and complicate the assignment of spikes to units based on their waveforms. We will address this problem by developing **spike sorting methods for long-duration and continual, long-duration and high-**

**channel-count recordings** (Figure 1, left box).

## **Aim 2: create real-time machine learning methods for intelligent experimentation**

In small-animal Neuroscience, most often statistical processing of neural time series is performed offline; i.e., experimental data is collected, saved to files, which are later statistically processed, with no runtime constraints. Most often all experimental data is processed at the same time; i.e., batch processing.

A new online statistical processing approach is now emerging in small-animal Neuroscience, where data is processed while it is being collected, and at the speed of data generation (?).

Online methods are well suited for NaLoDuCo experimentation. In experiments extending for weeks to months animals learn and adapt, their motivation and fatigue may fluctuate, and experimental conditions (e.g., lighting) may change. Offline batch processing algorithms cannot model this type of changing data. They assume stationary data whose statistical properties do not change across time. Differently, most online processing algorithms are robust to these changes. Also, NaLoDuCo experimentation is well suited for online methods, as the long-duration of these experiments provide a large amount of data to accurately fit expressive online methods.

We will **optimize methods developed for Aim 1 so that they can operate in real time**, and focus on the following two applications of these online methods (Figure 1, right box).

### **Intelligent neuromodulation**

Brain activity can be modulated optically, chemically and electrically (). Most commonly this modulation is done at fixed experimental times, or based on simple behavioral or neural observations.

We will guide optogenetic manipulations based on inferences from advanced machine learning methods. For example, a scientist may hypothesize that a peak in a neural latent variable, inferred from a prefrontal cortex population, signals the moment when mice decide to begin a foraging bout. To test this, she runs an online machine learning model to estimate latent variables from prefrontal cortex activity, predicting when this peak will occur. She then optogenetically inactivates the neural population at the forecasted

time. Because inactivation prevented the mouse from initiating a foraging bout, her hypothesis is supported.

### Intelligent experimental data storage

As the duration of NaLoDuCo experiments become longer, and the richness of the behavioral and neural recordings become larger, it will be unfeasible to store all raw data. We will be forced to intelligently decide, in real time, subsets of data to discard.

For instance, if we are recording videos from a mouse foraging in a large arena with ten high-resolution cameras, it would save considerable storage if at any time we only save videos from cameras capturing the mouse at that time. This could be done by tracking the position of the mouse in real time with probabilistic machine learning methods. Then, when the confidence of the tracking is high, we would only save videos of cameras capturing the mouse at the tracked position, but when the confidence is low, we would save all videos.

## 2 Approach

We have collected unprecedented NaLoDuCo datasets at the SWC. However, these very large datasets are of not much use without methods for visual exploration and data analysis. The next section describes methods we will disseminate for data analysis and Section 2.3 presents those for visualisation.

### 2.1 Offline analysis methods

We will disseminate an open-source library of advanced statistical methods adapted to **efficiently** process **non-stationary** recordings from NaLoDuCo experiments. We will benchmark the performance of these methods for processing behavioral and neural time series recorded from the SWC NaLoDuCo foraging experiments.

The initial list of methods in this library is given in Section ?? and includes regression, classification, supervised, unsupervised and reinforcement learning, state space models and artificial neural networks. Implementations of these methods will follow high software engineering standards, and will include detailed documentation, so that users can easily apply them to process

their own NaLoDuCo datasets. If a user needs a machine learning method to process NaLoDuCo recordings not included in the library, she could use implementations of similar algorithms included in the library, and their corresponding documentation, to adapt this algorithm to efficiently process her recordings. This library should be built for all and by all, and become essential to the rapidly expanding community of scientists performing NaLoDuCo experimentation in small animals around the world ().

In Neuroscience we don't have methods to characterise long-duration and continual time series, to learn from time series whose statistical properties fluctuate over time, to forecast time series of long horizons (e.g., hours, days, week or months. This library should find applicability beyond the realm of small-animal neuroscience and, for example, be useful to characterize long-duration and continual neural processes in human brain activity measured with subscalp EEG electrodes.

## 2.2 Initial list of methods

The first step in the analysis of NaLoDuCo foraging behavioral data is **tracking multiple body parts** in mice. For this we will use **deep learning** methods, as in (?). Next, we will use the previous tracking outputs to **infer mice kinematics** with **linear dynamical models**, as in (?). We will combine the tracking outputs with the kinematics inferences to **infer behavioral states** with **hidden Markov models**, as in (?). Further, we will **related kinematics and behavioral states to the probability of foraging events**, like leaving a patch, with **generalized linear models**, as in (), and **artificial neural networks**, as in (). The final step of the behavioral analysis will be to **infer mice policy from behavioral measures** with **inverse reinforcement learning**, as in ?.

The characterization of neural data will begin with the **estimation of latent variables models**, to reduced the dimensionality of multielectrode recordings of hundreds or even thousands of neurons to a small number of latent variables, using **latent variable models**, with linear (?) and non-linear (?) latent dynamics.. We will use the estimated latent variables as inputs to **infer neural states**, using **HMMs**, as in (). Next, we will **decode mice position** from hippocampal recordings, and **study replay** during long-duration **foraging**, with **point process decoders**, as in (?).

### 2.2.1 Challenges

Extracting meaning from long-duration and continual recordings opens challenges and opportunities that we will address and exploit in this project, as we describe in this and the next sections.

#### Non-stationarities

Conventional offline methods used to characterize neural time series assume that the statistical characteristics of the modeled data do not change with time (i.e., that the probability of the data is time invariant – stationarity). This assumption may be valid for shorter experiments. However, for long-duration experiments, where animals learn and adapt, where their motivation fluctuates, and their activity is modulated by circadian, ultradian and peridiem rhythms, this assumption may not hold. In nonstationary environments, a non-adaptive model trained under the false stationarity assumption is bound to become obsolete in time, and perform sub-optimally at best, or fail catastrophically at worst.

The field of adaptive signal processing develops algorithms to characterize non-stationary systems (?). In this field adaptations to specific algorithms have been developed to improve their performance in non-stationary environments.

For example, the recursive least-squares algorithm (?, Chapter 9) is an adaptation of the ordinary least square algorithm to perform **linear regression** with non-stationary data.

For non-linear regression using **artificial neural networks**, a very large number of strategies have been developed to address data non-stationarity. To mention a few, continual learning has introduced algorithms like Elastic Weight Consolidation (EWC ) and Learning Without Forgetting (LwF ) to allow models to adapt to changes over time without catastrophic forgetting. Also from this subfield is the Experience Replay (ER) algorithm that stores past data samples in a buffer and replays them alongside new data during training. A different type of strategy is used by ensemble methods (), which combine multiple models trained on different time windows to capture evolving data patterns.

Algorithms for **state-space models**, like the Kalman filter, perform well in some simpler non-stationary environments producing Gaussian data with varying mean and covariance. For more complex non-stationarities one can

use switching state-space models, like the switching linear dynamical system or the switching Hidden Markov model, or particle filters ().

In the machine learning literature, non-stationary methods are developed to address the problem of concept drift. It refers to the phenomenon where the statistical properties of data change over time, leading to a shift in the underlying relationships between variables. This presents a significant challenge for machine learning models, as assumptions made during training may become invalid, leading to degraded performance. Differently from adaptive signal processing, the methods developed to tackle concept drift are model agnostic and can be used with multiple machine learning models.

Concept drift can be categorized into different types based on the nature and rate of change:

**Sudden Drift:** A rapid and abrupt change in neural activity or behavioral response (e.g., a sudden change in neural firing patterns in response to a novel stimulus or a pharmacological intervention).

**Gradual Drift:** A slow transition in neural representations or behavior over time (e.g., gradual adaptation to a learned task, where neural activity patterns shift as the animal refines motor skills or cognitive performance).

**Incremental Drift:** A continuous and progressive evolution of neural activity or behavior (e.g., long-term potentiation (LTP) or synaptic plasticity leading to a steady, cumulative change in neural circuit function over time).

**Recurring or Seasonal Drift:** Periodic shifts in neural dynamics or behavior that return to previous states (e.g., circadian rhythm variations in neural firing rates, sleep/wake cycle-related changes in brain activity, or seasonal shifts in animal behavior such as hibernation or mating periods).

Detecting and adapting to concept drift is crucial for maintaining model performance. The primary approaches include:

**Statistical Tests and Distribution Monitoring:** Methods such as Kullback-Leibler divergence, Kolmogorov-Smirnov tests, and hypothesis testing can identify drift by comparing past and present distributions.



**Windowing Techniques:** Sliding windows, fixed or adaptive, monitor recent data points and adjust the model accordingly.

**Ensemble Methods:** Combining multiple models with different levels of adaptability to mitigate the effects of drift.

**Online Learning and Incremental Updates:** Continuously updating models with new data to reflect the latest trends.

**Drift-Aware Bayesian and Probabilistic Models:** Approaches such as Bayesian Online Learning allow probabilistic reasoning over changing data distributions.

**Unsupervised Drift Detection:** Methods such as clustering evolution, density estimation changes, and autoencoder-based monitoring can detect drift without labeled data.

### Long processing times for very large datasets

Neural and behavioral data analysis is most effective when computations can be performed quickly, ideally in real time. Very slow computations discourage data analysis, and hurts scientific discovery. The large dataset sizes generated by NaLoDuCo experimentation are an important challenge for fast data analysis.

To overcome this limitation, we will leverage distributed computing, a paradigm in which tasks and data are divided across multiple computers. Instead of relying on a single powerful machine, distributed computing accelerates processing by executing multiple parts of a computation in parallel.

We will develop parallel implementations of the of core machine learning algorithms for behavioral and neural data analysis (Section ??). These implementations will use Apache Spark<sup>1</sup> to parallelise pre-processing and feature extraction, and Ray<sup>2</sup> to parallelise machine learning and deep learning functionality.

### 2.2.2 Opportunities

**More expressive models** Our long-duration recordings, spanning weeks to months and generating hundreds of terabytes per experiment, will be

---

<sup>1</sup><https://spark.apache.org/>

<sup>2</sup><https://docs.ray.io/>

transformative for neuroscience, much like the advent of large-scale datasets in computer vision. Just as the creation of MNIST—and later, ImageNet—enabled the training of deeper neural networks, leading to unprecedented performance breakthroughs, our massive, high-resolution neural and behavioral datasets will allow the estimation of far more expressive models than previously possible. For instance, large NaLoDuCo datasets will allow to estimate latent variable models with highly nonlinear and expressive models of observation given latent variables using recognition parametrised models (?). By capturing neural dynamics over extended timescales, we may uncover novel insights into learning, memory, and long-term neural adaptations that remain inaccessible with conventional short-duration studies.

**Study very slow behavioral and neural rhythmic patterns** Continuously monitoring detailed behavioral and neural activity over weeks to months enables the study of slow rhythmic processes that extend beyond traditional circadian (24-hour) rhythms, including ultradian (hours), infradian (days to weeks), and even multi-month cycles. These long-duration fluctuations influence learning, memory consolidation, motivation, and cognitive function, yet they remain largely unexplored in controlled experiments. By capturing these dynamics, we can gain new insights into neural plasticity, attention, and mood regulation, as well as the progression of neurological disorders like Parkinson’s disease and depression, which exhibit slow symptom fluctuations.

**New neuromodulation opportunities** In traditional short-duration experiments, the effects of neuromodulation are tested immediately. In contrast, our long-duration experiments will enable repeated neuromodulation over extended periods and allow us to assess its impact over much longer timescales. For instance, in a mouse model of Alzheimer’s disease, we could apply optogenetic stimulation to the hippocampus for one hour per day over the course of a month and assess its impact on memory retention and synaptic plasticity in the following weeks. This approach could reveal whether intermittent neuromodulation promotes long-term neural circuit stability and delays cognitive decline.

### 2.2.3 Related research

**Neural data analysis methods from the Gatsby Unit** The Gatsby Unit has developed world-class neural data analysis methods for inferring latent variables using Gaussian processes (????), or variants of linear dynamical systems (??), or recognition parameterised models (?), for separating contributions of different factors to spiking activity using tensor decompositions (?), and for understanding the effects of neural perturbations (?), just to mention a few.

**Distributed computing for small animal Neuroscience** ? introduced a package<sup>3</sup> for analysing two-photon imaging records on distributed computing platforms. This has been a pioneering development by introducing distributing computing into neuroscience data analysis. However, it used short duration imaging recordings and it implemented simple data analysis methods.

**Continuous epilepsy monitoring** Continuous epilepsy monitoring is a transformative technology for diagnosing, understanding, and managing epilepsy. By capturing long-term brain activity, it provides crucial insights into seizure patterns and underlying neural dynamics.

A major advancement in this field is the development of implantable devices, such as the NeuroPace Responsive Neurostimulation (RNS) system, which continuously records electrocorticographic (ECoG) brain activity over extended periods (e.g., years). In addition to monitoring, RNS delivers targeted electrical stimulation in response to detected seizure precursors, significantly improving epilepsy treatment outcomes (?). However, RNS has two key limitations: (1) it is invasive, requiring surgical implantation, and (2) it can only store a limited amount of brain activity (typically a few hours) for later analysis, restricting its utility for studying long-term neural dynamics.

To overcome these challenges, researchers have developed subscalp ultra-long EEG recording technologies (?), which use electrodes implanted under the scalp but above the skull. This approach is less invasive than intracranial devices and offers continuous EEG monitoring, with the added benefit of streaming data directly to the cloud. This capability enables long-term characterization of brain activity.

---

<sup>3</sup><https://github.com/thunder-project/thunder>

Subscalp EEG monitoring has already provided valuable insights into epilepsy. For example, it has revealed that seizure susceptibility is often modulated by circadian and ultradian rhythms, with specific times of day associated with increased seizure likelihood. However, compared to the AEON platform, subscalp EEG provides lower-resolution measurements of both neural and behavioral activity. Additionally, the data modeling methods used in subscalp EEG research remain largely proprietary, as much of this technology is developed by private companies. Furthermore, current implementations are primarily designed for seizure detection and forecasting, rather than for broader investigations into long-term brain dynamics.

#### **2.2.4 Outputs**

Software repository containing implementations of machine learning algorithms relevant to NaLoDuCo experimentation, adapted to operate in non-stationary environments, and optimized to run on distributed computing environments (e.g., public clouds or institutional high-performance-computing clusters).

### **2.3 Visual Exploration**

#### **2.3.1 Outputs**

Visualisations for continuous behavioral and neural recording

Visualisations for epoched behavioral and neural recording

Visualisations for model outputs

Indexing system

### **2.4 Spike Sorting**

#### **2.4.1 Outputs**

Repository with implementations and benchmarking of offline spike sorting algorithms for long-duration recordings

Repository with implementations and benchmarking of online spike sorting algorithms

## **2.5 Dissemination**

### **2.5.1 Outputs**

- web application for the visualisation and analysis of SWC foraging NaLo-DuCo recordings in AWS
- documentation on how to build arenas and use the AEON platform and software
- documentation on how to use the online machine learning software

## **2.6 Online Machine Learning**

### **2.6.1 Outputs**

- Bonsai packages implementing real-time ML functionality for experimental control
- Documentation of these packages

## **2.7 Software and Infrastructure**